Horizontal Pod Autoscaler

Replace **<your-name>** with your **name** throughout the Lab.

1. SSH to the AWS Workstation

```
$ sudo su
# mkdir /home/devops/hpa
# cd /home/devops/hpa
# kubectl top nodes
```

```
root@ip-172-31-26-76:/home/devops# kubectl top nodes
NAME                                        CPU(cores)   CPU%   MEMORY(bytes)   MEMORY%
gke-hpa-demo-default-pool-795ab746-dgbs     72m          7%     770Mi           29%
gke-hpa-demo-pool-1-753df934-9fpt           50m          2%     521Mi           9%
root@ip-172-31-26-76:/home/devops#
```

2. Run & expose the Application

Create a new deployment with the below command.

```
# kubectl run hpa-demo-<your-name> --image=k8s.gcr.io/hpa-example
--requests=cpu=200m  --port=80
# kubectl expose deploy hpa-demo-<your-name> --type=NodePort
```

```
root@ip-172-31-26-76:/home/devops# kubectl run hpa-demo-albert --image=k8s.gcr.io/hpa-example --requests=cpu=200m  --port=80
kubectl run --generator=deployment/apps.v1 is DEPRECATED and will be removed in a future version. Use kubectl run --generator=run-pod/v1 or kubectl cr
eate instead.
deployment.apps/hpa-demo-albert created
root@ip-172-31-26-76:/home/devops#
```

```
root@ip-172-31-26-76:/home/devops# kubectl expose deploy hpa-demo-albert --type=NodePort
service/hpa-demo-albert exposed
root@ip-172-31-26-76:/home/devops#
```

3. Create Horizontal Pod Autoscaler

```
# kubectl autoscale deployment hpa-demo-<your-name> --cpu-percent=10
--min=1 --max=10
```

```
root@ip-172-31-26-76:/home/devops# kubectl autoscale deployment hpa-demo-albert --cpu-percent=10 --min=1 --max=10
horizontalpodautoscaler.autoscaling/hpa-demo-albert autoscaled
root@ip-172-31-26-76:/home/devops#
```

4. Please wait for 2-3 minutes before running the below command

```
# kubectl get hpa -w
```

```
root@ip-172-31-26-76:/home/devops# kubectl get hpa
NAME              REFERENCE                     TARGETS    MINPODS   MAXPODS   REPLICAS   AGE
hpa-demo-albert   Deployment/hpa-demo-albert   0%/10%     1         10        1          43s
root@ip-172-31-26-76:/home/devops#
```

Press Ctrl+c to exit.

7. Check the NODE where the HPA App has been deployed.

```
# kubectl get po -o wide
```

```
root@ip-172-31-26-76:/home/devops# kubectl get po -o wide
NAME                              READY   STATUS    RESTARTS   AGE   IP          NODE                                  NOMINATED NODE
hpa-demo-albert-6487b4997-kknzh   1/1     Running   0          3m    10.4.1.12   gke-hpa-demo-pool-1-753df934-9fpt   <none>
root@ip-172-31-26-76:/home/devops#
```

In this demo the POD is running on **gke-hpa-demo-pool-1-753df934-9fpt NODE.**

8. Run the below command to get the PUBLIC IP of the NODE where the POD is running.

```
# kubectl get nodes -o wide
```

```
root@ip-172-31-26-76:/home/devops# kubectl get nodes -o wide
NAME                                        STATUS   ROLES    AGE   VERSION          INTERNAL-IP   EXTERNAL-IP     OS-IMAGE                             K
ERNEL-VERSION   CONTAINER-RUNTIME
gke-hpa-demo-default-pool-795ab746-dgbs   Ready    <none>   1h    v1.11.7-gke.12   10.160.0.6    35.244.17.37    Container-Optimized OS from Google   4
.14.91+         docker://17.3.2
gke-hpa-demo-pool-1-753df934-9fpt         Ready    <none>   35m   v1.11.7-gke.12   10.160.0.7    35.244.57.29    Container-Optimized OS from Google   4
.14.91+         docker://17.3.2
root@ip-172-31-26-76:/home/devops#
```

Public IP for **gke-hpa-demo-pool-1-753df934-9fpt is 35.244.57.29.**

9. Check the NODEPORT on which the service is exposed.

```
# kubectl get svc hpa-demo-<your-name>
```

```
root@ip-172-31-26-76:/home/devops# kubectl get svc hpa-demo-albert
NAME              TYPE       CLUSTER-IP     EXTERNAL-IP   PORT(S)        AGE
hpa-demo-albert   NodePort   10.7.253.165   <none>        80:30514/TCP   7m
root@ip-172-31-26-76:/home/devops#
```

In this Example the hpa app is exposed is exposed on NODEPORT 30514

10. Run the below command to increase the Load.

```
while true; do wget -q -O- http://<NODE-Public-IP>:<NodePort>/ ; done
```

Where, NODE-Public-IP is the Public IP of the Node where the hpa-demo POD is deployed and NodePort is the Port on which the hpa-demo POD is exposed to, in this Example 35.244.57.29 is the Public IP of the NODE and 30514 is the NodePort

Example

```
while true; do wget -q -O- http://35.244.57.29:30514/ ; done
```

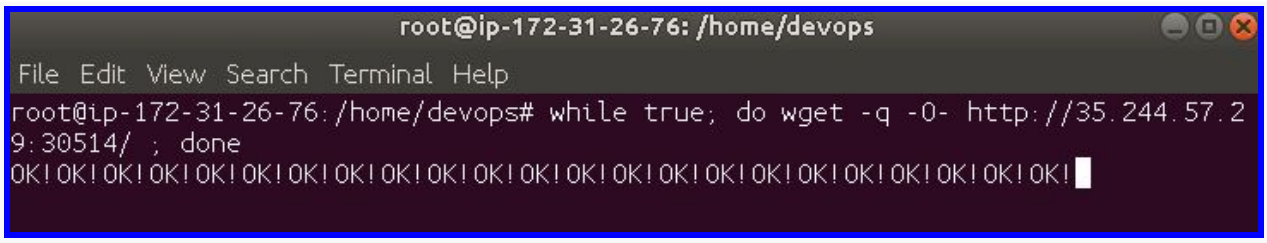

12. Launch TWO more terminals and SSH to your AWS Workstation from both the terminals.

13. On terminal two run the below command

```
$ sudo su
# watch -n 1 kubectl get hpa hpa-demo-<your-name>
```

14. On terminal three run the below command

```
$ sudo su
# watch -n 1 kubectl get pods
```

We can observe that the PODS have been horizontally scaled up due to the increased load on the third terminal, as shown in the below screenshot.

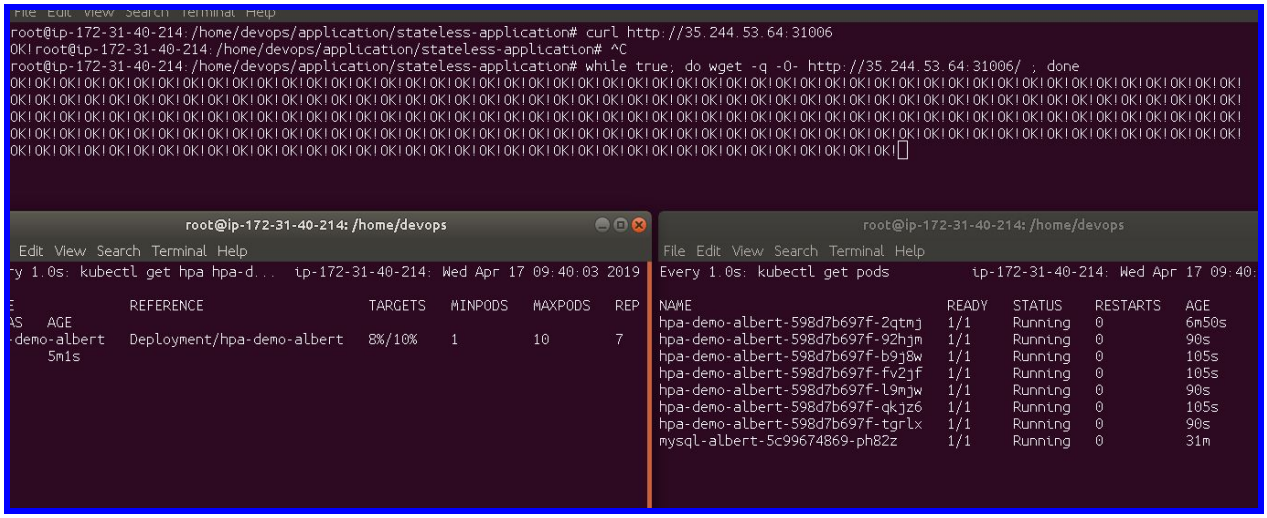

15. To Decrease the Load > Go back to the terminal ONE where we ran the command to increase the load and press CTRL+C to stop the load.

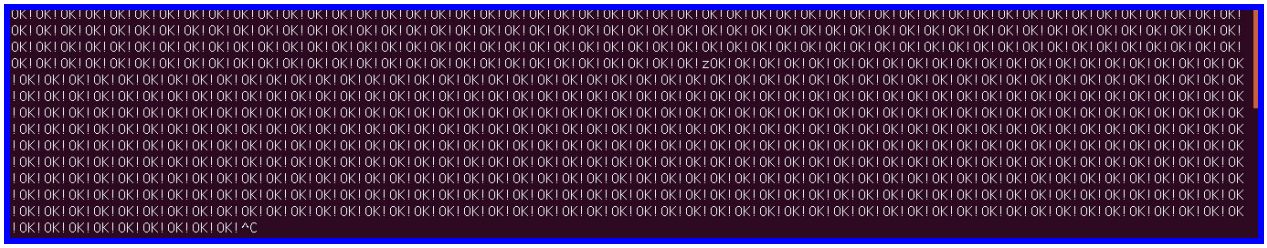

16. Wait for 5-7 minutes and go back to terminal TWO to check the output of the

```
# kubectl get pods
```

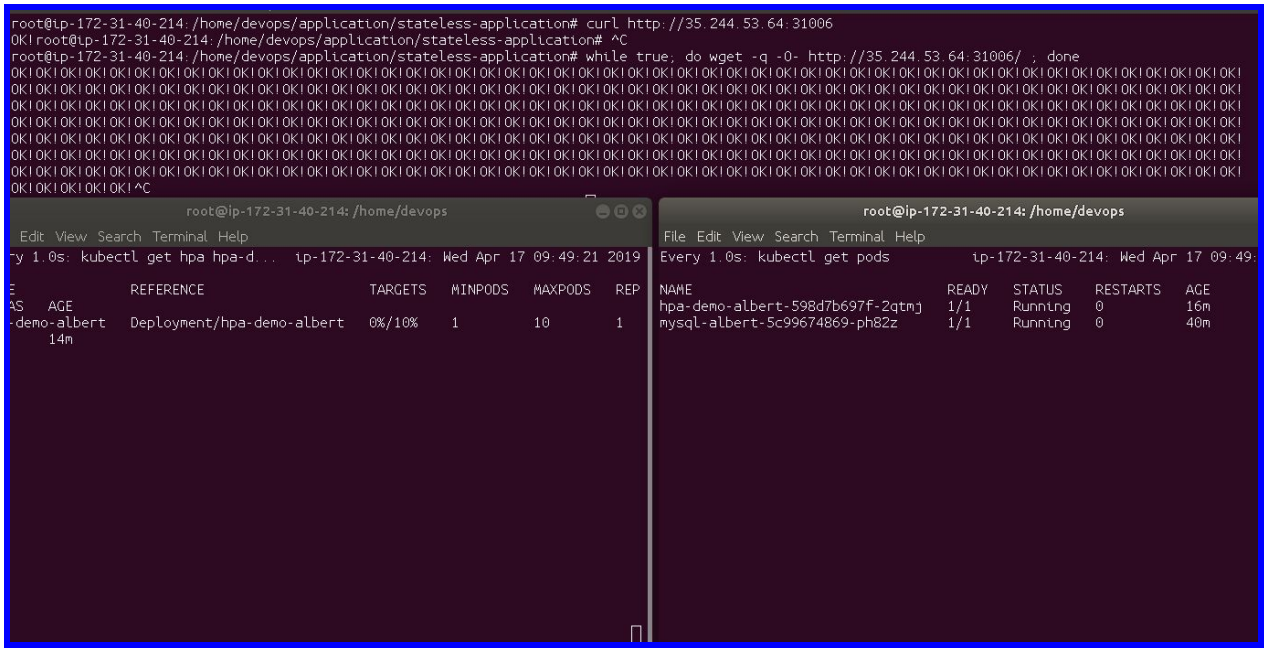

We can observe that the hpa-demo-albert app has scaled down to one again after the load is decreased.

```
# kubectl get hpa
```

Also, we can observe that the load has decreased to 0%.

19. On Terminal-1, Run the below commands to delete the Deployments and HPA

```
# kubectl delete deploy hpa-demo-<your-name>
# kubectl delete hpa hpa-demo-<your-name>
```